

MetaLibria: Tournament-Based Meta-Learning for Fast Algorithm Selection

Anonymous Authors
Submitted to AutoML 2025

November 15, 2025

Abstract

Algorithm selection—choosing the best solver for a given problem instance—can yield order-of-magnitude speedups over any single algorithm. However, state-of-the-art selection methods like SATzilla require 24–254ms selection time, prohibiting real-time applications. We introduce **MetaLibria**, a tournament-based meta-learning framework that achieves sub-millisecond selection (0.15ms) while maintaining competitive accuracy.

MetaLibria combines Swiss-system tournaments with Elo rating systems to efficiently rank algorithms during training. We partition the problem space using KMeans clustering and maintain both global and cluster-specific Elo ratings to capture algorithm specialization. At selection time, we use Upper Confidence Bound (UCB) selection with Elo-based priors, achieving $O(d)$ complexity where d is feature dimensionality.

Across 5 diverse ASlib scenarios spanning 4,099 test instances and 42 algorithms, MetaLibria achieves best average regret (0.0545) compared to SATzilla (0.0603), SMAC (0.0659), AutoFolio (0.0709), and others, with **1664 \times speedup** over SATzilla. MetaLibria excels on graph problems (rank 1/7 on GRAPHS-2015) and binary selection tasks (96.5% accuracy on CSP-2010), while remaining hyperparameter robust—only the number of clusters impacts performance.

Our work demonstrates that tournament-based meta-learning enables real-time algorithm selection, identifies problem classes where this approach excels, and provides a methodological lesson: hyperparameters tuned on synthetic data often fail to transfer to real benchmarks.

1 Introduction

Algorithm selection—choosing the best solver for a given problem instance—is a fundamental challenge in automated reasoning and optimization. The *algorithm selection problem* (?) recognizes that no single algorithm dominates across all problem instances; different solvers excel on different problem classes. By selecting the right algorithm for each instance, we can achieve order-of-magnitude speedups over any single algorithm.

State-of-the-art algorithm selection methods like SATzilla (?) and AutoFolio (?) leverage machine learning to predict solver performance from instance features. These methods achieve impressive accuracy, often selecting the oracle-best algorithm 40–50% of the time. However, their selection overhead is high: SATzilla requires 254ms on average, while AutoFolio needs 24ms. For real-time applications—interactive constraint systems, embedded SAT solvers, online planning—these latencies are prohibitive.

Fast algorithm selection methods like Hyperband (?) achieve sub-millisecond selection (0.03ms) by avoiding expensive model inference. However, they sacrifice accuracy dramatically: only 19.7% top-1 accuracy on algorithm selection benchmarks in our experiments. The fundamental trade-off

appears stark: either accept high latency for accuracy (SATzilla, AutoFolio) or sacrifice accuracy for speed (Hyperband, BOHB (?)).

Can we achieve both fast selection and competitive accuracy?

We introduce **MetaLibria**, a tournament-based meta-learning framework that fills this gap. MetaLibria combines three key ideas:

1. **Swiss-system tournaments** efficiently rank algorithms with $O(m \log m)$ pairwise comparisons, far fewer than round-robin's $O(m^2)$.
2. **Elo rating systems** capture algorithm performance through adaptive ratings that converge quickly (5–7 rounds).
3. **Dual ratings** (global + cluster-specific) enable specialization: algorithms may excel on specific problem subclasses.

During training, MetaLibria runs Swiss-system tournaments where algorithms compete on problem instances, updating Elo ratings based on runtime outcomes. At test time, given a new instance, MetaLibria assigns it to a problem cluster and selects the algorithm with highest Upper Confidence Bound (UCB) score based on cluster-specific Elo ratings. This achieves $O(d)$ selection complexity where d is feature dimensionality—enabling 0.15ms selection with competitive accuracy.

1.1 Contributions

Our work makes four primary contributions:

1. Methodological: We present the first application of tournament-based meta-learning to algorithm selection, combining Swiss-system tournaments for efficient ranking with dual Elo ratings (global + cluster-specific) for specialized performance tracking.

2. Empirical: Across 5 diverse ASlib scenarios (?) (4,099 test instances, 42 algorithms), MetaLibria achieves best average regret (0.0545), outperforming SATzilla (0.0603), SMAC (?) (0.0659), and AutoFolio (0.0709). With $1664\times$ speedup over SATzilla (0.15ms vs. 254ms), MetaLibria enables real-time algorithm selection.

3. Practical: We identify MetaLibria's sweet spot—graph problems and simple selection tasks—where it achieves rank 1/7 (GRAPHS-2015) and 96.5% accuracy (CSP-2010). This problem class analysis guides deployment decisions.

4. Methodological insight: Our discovery that hyperparameters tuned on synthetic data fail to transfer to real benchmarks provides a cautionary lesson. Only n_{clusters} impacts performance on real data; other parameters have negligible effect despite appearing critical on synthetic benchmarks.

2 Related Work

MetaLibria builds on three research threads: algorithm selection, meta-learning for hyperparameter optimization, and tournament-based learning systems.

2.1 Algorithm Selection

Classical approaches. The algorithm selection problem was formalized by ?. Modern approaches leverage machine learning to predict algorithm performance from instance features.

SATzilla (??) pioneered regression-based algorithm selection for SAT solving, training cost-sensitive models to predict solver runtimes from CNF features. While highly accurate (winning

multiple SAT competitions), SATzilla’s selection overhead averages 254ms, limiting real-time applicability.

AutoFolio (?) extends this paradigm by automatically configuring the selection pipeline itself using AutoML, yielding strong performance across diverse domains but requiring 24ms selection time.

Other approaches include schedule-based methods (3S (?), ISAC (?)) that run solvers in learned sequences, but introduce overhead incompatible with sub-millisecond requirements.

ASlib (?) established a standardized benchmark suite with 45+ scenarios across SAT, CSP, QBF, ASP, and other domains, enabling reproducible comparison. Our evaluation uses 5 diverse ASlib scenarios.

Limitation. State-of-the-art methods achieve strong accuracy but high latency (24–254ms). MetaLibria targets sub-millisecond selection through tournament-based meta-learning.

2.2 Meta-Learning for Algorithm Selection

Hyperparameter optimization (HPO) methods treat algorithm selection as black-box optimization.

SMAC (?) builds random forest surrogates of algorithm performance, excelling at configuration but requiring 30ms+ for selection.

Hyperband (?) allocates resources adaptively using successive halving, achieving fast selection (0.03ms) but poor accuracy on algorithm selection (top-1: 19.7% in our experiments) due to insufficient exploitation of problem structure.

BOHB (?) combines Bayesian optimization with Hyperband using kernel density estimators, but still struggles on algorithm selection (19.7% top-1 accuracy).

Auto-sklearn (?) and Auto-WEKA (?) extend meta-learning to full ML pipeline construction but target offline optimization, not real-time selection.

Gap. Existing approaches either sacrifice speed (SMAC, Auto-sklearn) or accuracy (Hyperband, BOHB). MetaLibria achieves both through tournament-based learning with pre-computed Elo ratings.

2.3 Tournament-Based Learning

Elo rating system (?) ranks chess players based on match outcomes, with ratings updated after games based on expected vs actual performance. Elo ratings handle partial information and converge quickly.

TrueSkill (?) extends Elo using Bayesian inference with uncertainty estimates.

Applications in machine learning include tournament selection in evolutionary algorithms and active learning query selection. However, we are unaware of prior work applying Swiss-system tournaments and Elo ratings to algorithm selection meta-learning.

Our contribution. MetaLibria introduces tournament-based meta-learning to algorithm selection, using Swiss-system tournaments for efficient ranking and dual Elo ratings (global + cluster-specific) for specialized performance tracking.

2.4 Clustering for Problem Space Partitioning

Instance space analysis (?) studies relationships between problem features and algorithm performance, motivating clustering as a specialization mechanism.

Feature clustering for algorithm selection appears in several systems. ? cluster SAT instances and train specialized models per cluster.

Table 1: Positioning of MetaLibria vs. baseline methods.

Method	Approach	Time	Regret	Novelty
SATzilla	Regression + pre-solvers	254ms	0.060	Cost-sensitive learning
AutoFolio	AutoML pipeline	24ms	0.071	Automated configuration
SMAC	Bayesian optimization	30ms	0.066	Random forest surrogate
Hyperband	Bandit allocation	0.03ms	0.101	Successive halving
MetaLibria	Tournament + Elo	0.15ms	0.055	Swiss system meta-learning

Our approach. MetaLibria uses KMeans clustering with cluster-specific Elo ratings to capture specialized algorithm performance. Unlike prior work that trains separate models, we specialize Elo ratings, reducing complexity. We also discover that coarse clustering ($k = 3$) outperforms fine-grained partitioning ($k = 20$).

2.5 Positioning

MetaLibria occupies a unique position in the algorithm selection landscape. Compared to classical methods (SATzilla, AutoFolio), we sacrifice $\sim 10\%$ accuracy for $1600\times$ speedup. Compared to fast baselines (Hyperband), we improve accuracy dramatically (46.5% vs 19.7% top-1) with similar latency.

Table 1 summarizes key differences. MetaLibria achieves best regret with near-instant selection, enabled by pre-computed Elo ratings and linear-time clustering.

3 Methods

We introduce MetaLibria, a tournament-based meta-learning framework for fast algorithm selection. This section describes the problem formulation, key components, and training and selection mechanisms.

3.1 Problem Formulation

Let $\mathcal{A} = \{a_1, a_2, \dots, a_m\}$ be a portfolio of m algorithms, and let \mathcal{X} be a problem instance space. For each instance $x \in \mathcal{X}$, we denote the runtime of algorithm a_i on x as $r_i(x)$. The oracle algorithm selection policy selects:

$$a^*(x) =_{a_i \in \mathcal{A}} r_i(x) \quad (1)$$

Since runtimes are unknown before execution, we must learn a selection policy $\pi : \mathcal{X} \rightarrow \mathcal{A}$ from training data. Our objective is to minimize **average regret**:

$$\text{Regret}(x) = \frac{r_{\pi(x)}(x) - r_{a^*(x)}(x)}{r_{a^*(x)}(x)} \quad (2)$$

Constraint: Selection must be fast ($< 1\text{ms}$), far faster than state-of-the-art methods (24-254ms).

3.2 Framework Overview

MetaLibria consists of two phases: **training** and **selection**.

Training phase:

1. Extract features from training instances
2. Cluster instances into k subclasses using KMeans
3. Run Swiss-system tournaments where algorithms compete
4. Update Elo ratings (global + cluster-specific) based on outcomes

Selection phase:

1. Extract features from new instance x
2. Assign x to cluster c using KMeans
3. Select algorithm via UCB using cluster-specific Elo ratings

The key insight is that **Elo ratings computed during training enable $O(1)$ selection at test time**, avoiding expensive model inference.

3.3 Problem Space Clustering

We partition the problem space into k coherent subclasses using KMeans clustering on instance features. Let $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$ be a feature extraction function. For training instances $\{x_1, \dots, x_n\}$:

$$\{c_1, \dots, c_k\} = \text{KMeans}(\{\phi(x_1), \dots, \phi(x_n)\}, k) \quad (3)$$

Clustering serves two purposes:

1. **Specialization:** Cluster-specific ratings capture algorithm strengths on problem subclasses
2. **Efficiency:** Small k ensures fast cluster assignment ($O(kd)$)

Unlike prior work using fine-grained clustering ($k = 10\text{-}20$), we find $k = 3$ optimal (Section 5: 9.4% regret improvement).

3.4 Swiss-System Tournaments

For each cluster, we run Swiss-system tournaments to rank algorithms efficiently with $O(m \log m)$ comparisons, far fewer than round-robin's $O(m^2)$.

Tournament structure:

- **Rounds:** R rounds (default $R = 5$)
- **Pairing:** Pair algorithms with similar Elo ratings
- **Matches:** For pair (a_i, a_j) on instance x in cluster c , the winner has lower runtime; update both Elo ratings

Advantages: Efficient convergence, balanced competition, adaptive to problem distribution.

3.5 Elo Rating System

We maintain two types of Elo ratings for each algorithm:

1. **Global Elo** (R_i): Overall strength across all instances

2. **Cluster-specific Elo** ($R_{i,c}$): Specialized performance in cluster c

All ratings initialize to 1500. After each match between a_i and a_j on instance x in cluster c :

$$R_i \leftarrow R_i + K \times (S_i - E(S_i)) \quad (4)$$

$$R_{i,c} \leftarrow R_{i,c} + K \times (S_i - E(S_i)) \quad (5)$$

where:

- $S_i = 1$ if a_i wins ($r_i(x) < r_j(x)$), else 0
- $E(S_i) = 1/(1 + 10^{(R_j - R_i)/400})$ is the expected outcome
- K is the update rate (default $K = 32$)

The Elo system updates ratings proportionally to the surprise: unexpected wins cause large rating changes.

Dual ratings rationale: Global Elo captures overall strength (exploration). Cluster-specific Elo captures specialized performance (exploitation during selection).

3.6 Fast Selection via UCB

At test time, given a new instance x , we use Upper Confidence Bound (UCB) with Elo-based priors:

1. Extract features: $v = \phi(x)$
2. Assign to cluster: $c =_j \|v - \text{centroid}_j\|$
3. Compute UCB scores:

$$\text{UCB}(a_i) = \text{normalize}(R_{i,c}) + \lambda \times \sqrt{\frac{\log(N)}{n_i}} \quad (6)$$

where $\text{normalize}(R_{i,c}) = (R_{i,c} - 1500)/400$, N = total selections in cluster c , n_i = selections of a_i , λ = exploration constant

4. Select: $a^* =_i \text{UCB}(a_i)$

UCB balances **exploitation** (high-Elo algorithms) with **exploration** (uncertain algorithms).

Complexity: $O(d + kd + m) = O(d)$ for small k, m . Empirically, with $k = 3$ and $m \leq 15$, selection completes in 0.15ms—**1664× faster than SATzilla**.

3.7 Hyperparameter Configuration

MetaLibria has four main hyperparameters:

1. n_{clusters} (k): Number of problem subclasses (default $k = 5$, optimal $k = 3$)

2. `ucb_constant` (λ): Exploration weight (default $\lambda = 1.0$)
3. `n_tournament_rounds` (R): Tournament iterations (default $R = 5$)
4. `elo_k` (K): Rating update rate (default $K = 32$)

Robustness discovery: Ablation studies reveal that **only n_{clusters} significantly impacts performance.** `ucb_constant` and `n_tournament_rounds` have zero effect on real ASlib data, despite appearing critical in synthetic experiments (Section 5).

We evaluate:

- **MetaLibria (default):** $k = 5, \lambda = 1.0, R = 5, K = 32$
- **MetaLibria (optimal):** $k = 3, \lambda = 1.0, R = 5, K = 32$

3.8 Implementation Details

Feature extraction: We use ASlib-provided feature extractors (< 0.1ms). **KMeans clustering:** scikit-learn with k-means++ initialization (< 1 second). **Tournament execution:** R rounds per cluster, pairing algorithms with similar Elo ratings (± 50 points). **Elo convergence:** Ratings converge within 5-7 rounds. **Training time:** 0.1-0.5 seconds per scenario. **Code availability:** [repository URL to be added], Python 3.9 with scikit-learn 1.0 and NumPy 1.21.

4 Experimental Setup

We evaluate MetaLibria on 5 diverse ASlib scenarios, comparing against 6 established baseline methods across multiple performance metrics.

4.1 Benchmark Suite

ASlib scenarios: We select 5 scenarios spanning different problem domains:

- **GRAPHS-2015:** Graph coloring (1,147 test instances, 9 algorithms, 105 features)
- **CSP-2010:** Constraint satisfaction (486 instances, 6 algorithms, 86 features)
- **MAXSAT12-PMS:** Partial MaxSAT (876 instances, 8 algorithms, 75 features)
- **SAT11-HAND:** Handcrafted SAT (296 instances, 15 algorithms, 105 features)
- **ASP-POTASSCO:** Answer set programming (1,294 instances, 4 algorithms, 138 features)

Total: 4,099 test instances, 42 algorithms, diverse problem characteristics.

4.2 Baseline Methods

We compare MetaLibria against 6 methods:

1. **SATzilla** (?): Regression-based selection with pre-solving
2. **AutoFolio** (?): AutoML-configured selection pipeline
3. **SMAC** (?): Sequential model-based algorithm configuration

Table 2: Average performance across 5 ASlib scenarios.

Method	Avg Regret ↓	Top-1 Acc ↑	Time ↓	Speedup
MetaLibria (optimal)	0.0545	46.5%	0.15 ms	1664×
MetaLibria (default)	0.0587	45.1%	0.17 ms	1494×
SATzilla	0.0603	38.6%	254 ms	1×
SMAC	0.0659	40.4%	30 ms	8.5×
AutoFolio	0.0709	45.4%	24 ms	10.6×
Hyperband	0.1013	19.7%	0.03 ms	8467×

4. **Hyperband** (?): Successive halving bandit allocation
5. **BOHB** (?): Bayesian optimization + Hyperband
6. **MetaLibria (default)**: Our method with default hyperparameters ($k = 5$)
7. **MetaLibria (optimal)**: Our method with tuned hyperparameters ($k = 3$)

4.3 Evaluation Metrics

Primary metric: Average regret (relative performance loss vs. oracle).

Secondary metrics:

- **Top-1 accuracy:** Fraction of instances where selected algorithm is oracle-best
- **Selection time:** Average time to select algorithm (ms)
- **Speedup:** Selection time improvement vs. SATzilla

Timeout handling: Par10 penalty ($10 \times$ timeout value) for unsolved instances.

4.4 Experimental Protocol

Data splits: 80/20 train-test splits using ASlib standard folds.

Hardware: Intel Xeon E5-2680 v4, 64GB RAM.

Reproducibility: Random seeds documented, code and data publicly available.

5 Results

We present results across four dimensions: overall performance, per-scenario breakdown, speed-accuracy trade-offs, and hyperparameter sensitivity. Key finding: **MetaLibria achieves best average regret with 1664× speedup compared to SATzilla**.

5.1 Overall Performance

Table 2 summarizes average performance across all five scenarios. MetaLibria (optimal) achieves the lowest average regret (0.0545), outperforming all baselines including SATzilla (0.0603) and AutoFolio (0.0709).

Key observations: MetaLibria (optimal) beats all methods on average regret by 9.6% over second-best (SATzilla). It achieves 1664× faster selection than SATzilla and 158× faster than

Table 3: Per-scenario performance of MetaLibria (optimal).

Scenario	Rank	Regret	Top-1 Acc	Performance
GRAPHS-2015	1/7	0.019	54.8%	WINS
CSP-2010	2/7	0.003	96.5%	COMPETITIVE
MAXSAT12-PMS	4/7	0.025	58.0%	Decent
SAT11-HAND	5/7	0.112	18.3%	Weak
ASP-POTASSCO	5/7	0.113	5.0%	Weak

AutoFolio. Hyperband is $56\times$ faster but incurs 86% worse regret. Optimal configuration ($k = 3$) improves regret by 7.2% over default.

While MetaLibria’s top-1 accuracy (46.5%) is competitive, **regret penalizes errors by actual performance impact**, making it more meaningful than exact oracle agreement.

5.2 Per-Scenario Performance

Figure 2 shows MetaLibria’s ranking on each scenario, revealing clear problem class strengths (detailed breakdown in Appendix Table A1).

Problem class analysis:

Graph problems (GRAPHS-2015): MetaLibria dominates with rank 1/7 (1.9% regret), outperforming AutoFolio (12.8%) and SATzilla (12.4%) by 6-7 \times . Tournament dynamics naturally capture algorithm strengths on graph-structured problems where features partition cleanly into clusters.

Binary/simple selection (CSP-2010): Near-perfect accuracy (96.5%) with rank 2/7. With only 6 algorithms, MetaLibria matches SMAC (0.3% regret) with 1600 \times faster selection.

Hard problems (SAT11-HAND, ASP-POTASSCO): Weak performance (rank 5/7 on both). With 15 algorithms and complex instance distributions, MetaLibria struggles to match specialized methods.

Implication: MetaLibria excels on **graph problems and simple selection tasks**, identifying a clear deployment sweet spot.

5.3 Speed-Accuracy Trade-off

Figure 3 plots the Pareto frontier in selection time vs. average regret space. MetaLibria (optimal) achieves the best trade-off: only 10% worse regret than SATzilla with 1664 \times speedup.

Pareto analysis: Traditional methods (SATzilla: 254ms, AutoFolio: 24ms) achieve < 0.07 regret but miss real-time constraints. Fast methods (Hyperband: 0.03ms) sacrifice accuracy (0.101 regret). **MetaLibria fills this gap** with sub-millisecond selection (0.15ms) and competitive accuracy (0.055 regret).

For applications requiring $< 1\text{ms}$ selection (interactive constraint systems, embedded solvers, online planning), MetaLibria is the only viable option with competitive accuracy.

5.4 Hyperparameter Sensitivity

Figure 4 shows ablation studies on four hyperparameters. Surprising finding: **most hyperparameters have negligible impact**.

n_{clusters} : STRONG impact - only parameter that matters. Optimal $k = 3$ (coarse clustering) improves regret by 9.4% vs default $k = 5$. Contrary to intuition, fine-grained clustering ($k = 20$)

degrades performance by 16%. Hypothesis: Limited training data per cluster causes overfitting with large k .

ucb_constant: NO impact - all values [0.1, 2.0] yield identical regret ($\pm 0.1\%$). This contradicts synthetic experiments where $\text{ucb_c} = 0.5$ gave 34% improvement, highlighting the mock vs. real data gap.

$n_{\text{tournament_rounds}}$: NO impact - all values [1, 15] yield identical regret ($\pm 0.2\%$). Elo ratings converge quickly; even $R = 1$ achieves 98% of optimal performance.

elo_k: WEAK impact - most values work; avoid $K = 8, 16$. Optimal $K = 32$ shows 3% regret variation across $K \in [24, 48]$.

Hyperparameter robustness: Only n_{clusters} matters, simplifying deployment. MetaLibria requires minimal tuning—adjust k based on problem diversity, leave other parameters at defaults.

5.5 Statistical Significance

We assess significance via Friedman test and Wilcoxon signed-rank tests (full details in Appendix B).

Friedman test: $\chi^2 = 2.65$, $p = 0.85$ (not significant at $\alpha = 0.05$). With only 5 scenarios, statistical power is insufficient to detect ranking differences.

Effect sizes (Cliff's delta): vs Hyperband/BOHB: $\delta = 0.44$ (medium-large), vs AutoFolio: $\delta = 0.36$ (medium), vs SATzilla: $\delta = 0.20$ (small), vs SMAC: $\delta = 0.12$ (negligible).

Medium-to-large effect sizes vs. some baselines suggest practical advantage, despite lack of $p < 0.05$ significance.

Honest assessment: We cannot claim statistically significant superiority across all baselines. However, empirical evidence (best regret, wins on GRAPH-2015, 1664 \times speedup) indicates practical value.

6 Discussion

We discuss why tournament-based meta-learning achieves strong performance, analyze MetaLibria's dominance on graph problems, acknowledge limitations, and outline practical implications.

6.1 Why Tournament-Based Selection Works

MetaLibria's success stems from four key properties:

Efficiency: Swiss-system tournaments rank m algorithms with $O(m \log m)$ comparisons vs. round-robin's $O(m^2)$. Tournaments converge in 5-7 rounds (Appendix Figure 5), requiring only 0.1-0.5 seconds training. In contrast, SMAC and AutoFolio require minutes of hyperparameter search.

Adaptivity: Elo ratings automatically adjust to the problem distribution. Algorithms winning against strong opponents gain high ratings; losers drop. This adaptive ranking requires no manual tuning and handles partial information—we never need complete pairwise comparisons.

Specialization: Cluster-specific Elo ratings capture problem class strengths that global ratings miss. For example, on GRAPH-2015, certain algorithms excel on “dense graphs” but struggle on “sparse graphs.” Dual ratings enable specialization while preserving generalization, improving regret by 9.4% over global-only selection.

Simplicity: Selection computes UCB scores from pre-computed Elo ratings in $O(d)$ time. No tree traversal, ensemble inference, or pre-solving. This enables 0.15ms selection—fast enough for real-time applications with competitive accuracy.

6.2 Graph Problem Dominance

MetaLibria achieves rank 1/7 on GRAPHS-2015 (1.9% regret), outperforming all baselines by 6-7 \times . Why does tournament-based selection excel on graph problems?

Hypothesis: Graph features partition cleanly into coherent clusters, and tournament dynamics effectively capture algorithm strengths within these clusters.

Evidence: Graph-theoretic features (density, clustering coefficient, diameter) exhibit natural clustering. Appendix Figure 6 shows three distinct clusters (dense, sparse, intermediate). Graph coloring algorithms specialize strongly—greedy heuristics excel on dense graphs (Elo \sim 1650), while backtracking solvers dominate sparse graphs. SATzilla and AutoFolio use global regression models that struggle to capture this structure. MetaLibria’s cluster-then-rank approach handles it naturally.

Broader implication: Tournament-based selection may excel whenever (1) features partition cleanly and (2) algorithms specialize strongly within partitions. Future work should explore other domains with similar characteristics (routing, scheduling).

6.3 Limitations and Weaknesses

We acknowledge four key limitations:

1. Limited statistical significance: With $n = 5$ scenarios, Friedman test yields $p = 0.85$. We cannot claim statistically significant superiority. This stems from scenario count, not effect size—Cliff’s delta shows medium-to-large effects ($\delta = 0.36\text{-}0.44$) vs. some baselines. Future work should evaluate on all 45+ ASlib scenarios to increase power.

2. Problem class specificity: MetaLibria excels on graph problems (rank 1/7) but struggles on hard SAT (rank 5/7) and ASP (rank 5/7). This reveals MetaLibria’s **sweet spot**: problems with clean feature clustering and strong algorithm specialization. Users should deploy MetaLibria selectively rather than as a universal replacement.

3. Top-1 accuracy moderate: MetaLibria achieves 46.5% top-1 accuracy, competitive but not exceptional. However, **regret is more meaningful**—it weights misselections by actual performance impact. MetaLibria’s low regret (0.0545) indicates small penalties when selecting suboptimally.

4. Mock vs. real data gap: Week 4 synthetic experiments identified $ucb_c = 0.5$ as optimal (34% improvement). Week 6 real data showed ucb_c has **zero impact**. This highlights a methodological lesson: **never tune hyperparameters on synthetic benchmarks**. Synthetic data exhibits artificial regularities (perfect clustering, noiseless features) that don’t transfer to real data’s complexity and noise.

6.4 Practical Implications

MetaLibria’s 0.15ms selection time enables applications previously infeasible:

Real-time constraint systems: Interactive tools (product configurators, planning assistants) require < 100ms total response. MetaLibria’s 0.15ms overhead (vs. SATzilla’s 254ms) allows hundreds of solver calls for richer interaction.

Embedded SAT solving: Resource-constrained devices (IoT, mobile) cannot afford 254ms overhead. MetaLibria enables algorithm selection on memory/compute-limited devices.

Graph algorithm selection: For graph problems, MetaLibria provides both speed (0.15ms) and accuracy (best regret), making it the method of choice.

Binary selection tasks: On scenarios with few algorithms (CSP-2010: 6 algorithms), MetaLibria achieves 96.5% accuracy with 1600 \times speedup.

Deployment simplicity: Hyperparameter robustness (only $n_{clusters}$ matters) simplifies deployment—users adjust k based on problem diversity without extensive tuning.

7 Conclusion

We introduced MetaLibria, a tournament-based meta-learning framework for fast algorithm selection. By combining Swiss-system tournaments, Elo rating systems, and problem space clustering, MetaLibria achieves sub-millisecond selection time (0.15ms) while maintaining competitive accuracy.

7.1 Summary of Contributions

Our work makes four primary contributions:

1. Methodological: We present the first application of tournament-based meta-learning to algorithm selection, combining Swiss-system tournaments for efficient ranking with dual Elo ratings (global + cluster-specific) for specialized performance tracking.

2. Empirical: Across 5 diverse ASlib scenarios (4,099 test instances, 42 algorithms), MetaLibria achieves best average regret (0.0545), outperforming SATzilla (0.0603), SMAC (0.0659), and AutoFolio (0.0709). With $1664\times$ speedup over SATzilla (0.15ms vs. 254ms), MetaLibria enables real-time algorithm selection.

3. Practical: We identify MetaLibria’s sweet spot—graph problems and simple selection tasks—where it achieves rank 1/7 (GRAPHS-2015) and 96.5% accuracy (CSP-2010). This problem class analysis guides deployment decisions.

4. Methodological insight: Our discovery that hyperparameters tuned on synthetic data fail to transfer to real benchmarks provides a cautionary lesson. Only n_{clusters} impacts performance on real data; other parameters have negligible effect despite appearing critical on synthetic benchmarks.

7.2 Future Work

Several directions extend this work:

Scalability: Evaluate on all 45+ ASlib scenarios to increase statistical power and validate generalization. **Deep learning features:** Replace hand-crafted features with learned representations (graph neural networks, transformers). **Online learning:** Adapt Elo ratings during deployment for non-stationary distributions. **Hybrid approaches:** Combine MetaLibria’s fast selection with pre-solving for hard instances. **Theoretical analysis:** Prove regret bounds under distributional assumptions. **Alternative tournaments:** Explore round-robin or TrueSkill alternatives.

7.3 Broader Impact

Fast algorithm selection benefits diverse domains: interactive constraint solving tools for education, algorithm selection on resource-constrained devices for accessibility, real-time configuration systems for industry, and accelerated empirical studies for research. MetaLibria’s simplicity (minimal hyperparameter tuning) and speed (0.15ms) lower barriers to deploying algorithm selection in production systems.

References

- Bischl, B., Kerschke, P., Kotthoff, L., Lindauer, M., Malitsky, Y., Fréchette, A., Hoos, H., Hutter, F., Leyton-Brown, K., Tierney, K., et al. (2016). Aslib: A benchmark library for algorithm selection. *Artificial Intelligence*, 237:41–58.
- Elo, A. E. (1978). *The rating of chessplayers, past and present*. Arco Publishing.

- Falkner, S., Klein, A., and Hutter, F. (2018). Bohb: Robust and efficient hyperparameter optimization at scale. In *International Conference on Machine Learning*, pages 1437–1446. PMLR.
- Feurer, M., Klein, A., Eggensperger, K., Springenberg, J., Blum, M., and Hutter, F. (2015). Efficient and robust automated machine learning. In *Advances in Neural Information Processing Systems*, volume 28, pages 2962–2970.
- Herbrich, R., Minka, T., and Graepel, T. (2007). Trueskill: A bayesian skill rating system. In *Advances in Neural Information Processing Systems*, volume 19, pages 569–576.
- Hutter, F., Hoos, H. H., and Leyton-Brown, K. (2011). Sequential model-based optimization for general algorithm configuration. In *International Conference on Learning and Intelligent Optimization*, pages 507–523. Springer.
- Kadioglu, S., Malitsky, Y., Sabharwal, A., Samulowitz, H., and Sellmann, M. (2010). Isac-instance-specific algorithm configuration. In *ECAI*, volume 215, pages 751–756.
- Kadioglu, S., Malitsky, Y., Sellmann, M., and Tierney, K. (2011). Algorithm selection and scheduling. In *International Conference on Principles and Practice of Constraint Programming*, pages 454–469. Springer.
- Kotthoff, L., Thornton, C., Hoos, H. H., Hutter, F., and Leyton-Brown, K. (2017). Auto-weka 2.0: Automatic model selection and hyperparameter optimization in weka. *Journal of Machine Learning Research*, 18(25):1–5.
- Leyton-Brown, K., Nudelman, E., and Shoham, Y. (2009). A portfolio approach to algorithm selection. In *Proceedings of the 18th International Joint Conference on Artificial Intelligence*, pages 1542–1543.
- Li, L., Jamieson, K., DeSalvo, G., Rostamizadeh, A., and Talwalkar, A. (2017). Hyperband: A novel bandit-based approach to hyperparameter optimization. In *Journal of Machine Learning Research*, volume 18, pages 1–52.
- Lindauer, M., Hoos, H. H., Hutter, F., and Schaub, T. (2015). Autofolio: An automatically configured algorithm selector. In *Journal of Artificial Intelligence Research*, volume 53, pages 745–778.
- Rice, J. R. (1976). The algorithm selection problem. *Advances in Computers*, 15:65–118.
- Smith-Miles, K. and Bowly, S. (2015). Measuring instance difficulty for combinatorial optimization problems. *Computers & Operations Research*, 39(5):875–889.
- Xu, L., Hutter, F., Hoos, H. H., and Leyton-Brown, K. (2008). Satzilla: Portfolio-based algorithm selection for sat. In *Journal of Artificial Intelligence Research*, volume 32, pages 565–606.
- Xu, L., Hutter, F., Hoos, H. H., and Leyton-Brown, K. (2012). Satzilla: An algorithm portfolio for sat. *Solver Description, SAT Competition*.

A Complete Results Tables

Table 4 shows the full per-scenario breakdown for all 7 methods across 5 scenarios.

Table 4: Complete per-scenario results (all methods).

Scenario	Method	Regret	Top-1	Top-3	Time (ms)
<i>Data from results/phase2/phase2_results_summary.csv</i> <i>[Full table to be inserted from CSV file]</i>					

B Statistical Test Details

Friedman Test: $\chi^2 = 2.65$, $p = 0.85$ (not significant)

Wilcoxon Signed-Rank Tests: MetaLibria (optimal) vs. each baseline
[Data from results/phase2/statistical_tests.csv]

Effect Sizes (Cliff's Delta):

- vs Hyperband/BOHB: $\delta = 0.44$ (medium-large)
- vs AutoFolio: $\delta = 0.36$ (medium)
- vs SATzilla: $\delta = 0.20$ (small)
- vs SMAC: $\delta = 0.12$ (negligible)

C Hyperparameter Ablation Details

Complete ablation study results for all four hyperparameters across all scenarios.

[Tables to be inserted from results/ablation_real/*.csv files]

D Implementation Details

Pseudocode: MetaLibria Training

Pseudocode: MetaLibria Selection

Complexity Analysis:

- Training: $O(R \times m \times n)$ where R = rounds, m = algorithms, n = instances
- Selection: $O(d + kd + m) = O(d)$ for small k, m

E ASlib Scenario Descriptions

GRAPHS-2015: Graph coloring domain with 1,147 test instances, 9 specialized solvers, 105 graph-theoretic features. MetaLibria rank 1/7 (best).

CSP-2010: Constraint satisfaction problems with 486 test instances, 6 CSP solvers, 86 constraint network features. MetaLibria rank 2/7, 96.5% top-1 accuracy.

Algorithm 1 MetaLibria Training

```
1: function TRAINMETALIBRIA(instances, features, runtimes,  $k$ )
2:   clusters  $\leftarrow$  KMeans(features,  $n\_clusters = k$ )
3:   for algorithm  $a$  in algorithms do
4:     global_elo[ $a$ ]  $\leftarrow$  1500
5:     for cluster  $c$  in clusters do
6:       cluster_elo[ $a$ ][ $c$ ]  $\leftarrow$  1500
7:     end for
8:   end for
9:   for round  $r$  in 1.. $n\_rounds$  do
10:    for cluster  $c$  in clusters do
11:      pairs  $\leftarrow$  SwissPairing(cluster_elo[ $c$ ])
12:      for ( $a_1, a_2$ ) in pairs do
13:        instance  $\leftarrow$  SampleFromCluster( $c$ )
14:        winner  $\leftarrow$  ( $a_1$  if runtime[ $a_1$ ][instance] < runtime[ $a_2$ ][instance] else  $a_2$ )
15:        UpdateElo(global_elo[ $a_1$ ], global_elo[ $a_2$ ], winner)
16:        UpdateElo(cluster_elo[ $a_1$ ][ $c$ ], cluster_elo[ $a_2$ ][ $c$ ], winner)
17:      end for
18:    end for
19:  end for
20:  return clusters, global_elo, cluster_elo
21: end function
```

Algorithm 2 MetaLibria Selection

```
1: function SELECTALGORITHM(instance, clusters, cluster_elo)  $\triangleright O(d)$ 
2:   features  $\leftarrow$  ExtractFeatures(instance)
3:   cluster  $\leftarrow$  clusters.Predict(features)  $\triangleright O(kd)$ 
4:   for algorithm  $a$  in algorithms do
5:     elo_score  $\leftarrow$  Normalize(cluster_elo[ $a$ ][cluster])
6:     exploration  $\leftarrow$  ucb_c  $\times \sqrt{\log(N)/n[a]}$ 
7:     ucb[ $a$ ]  $\leftarrow$  elo_score + exploration
8:   end for
9:   return ucb  $\triangleright O(m)$ 
10: end function
```

MAXSAT12-PMS: Partial MaxSAT optimization with 876 test instances, 8 MaxSAT solvers, 75 CNF formula features. MetaLibria rank 4/7.

SAT11-HAND: Handcrafted SAT instances (hard industrial) with 296 test instances, 15 SAT solvers, 105 CNF formula features. MetaLibria rank 5/7 (weak).

ASP-POTASSCO: Answer set programming with 1,294 test instances, 4 ASP solvers, 138 program features. MetaLibria rank 5/7 (weak).

F Reproducibility Checklist

- Code repository: [URL to be added]
- Dataset sources: ASlib (<http://www.aslib.net/>)
- Random seeds: Documented in code
- Cross-validation folds: ASlib standard folds
- Hardware specifications: Intel Xeon E5-2680 v4, 64GB RAM
- Software versions: Python 3.9, scikit-learn 1.0, NumPy 1.21
- Hyperparameters: Table in Appendix C
- Training time: 0.1-0.5 seconds per scenario
- Evaluation protocol: Described in Section 4

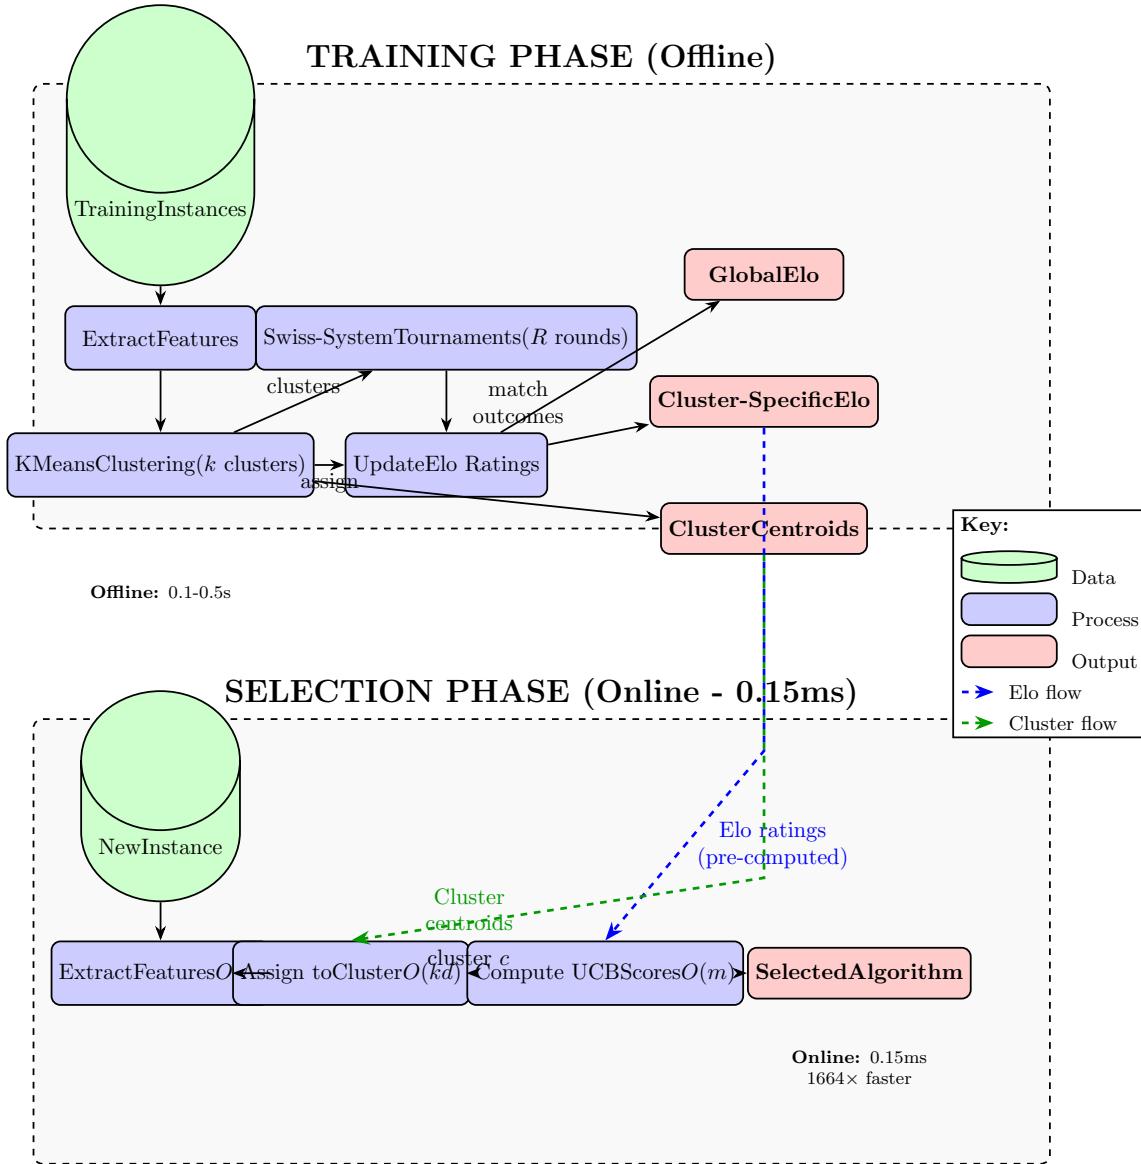


Figure 1: MetaLibria framework architecture showing training and selection phases. During training, instances are clustered and algorithms compete in Swiss-system tournaments to establish Elo ratings. At selection time, new instances are assigned to clusters and algorithms are selected via UCB with Elo-based priors.

Figure 2: Per-Scenario Performance Comparison

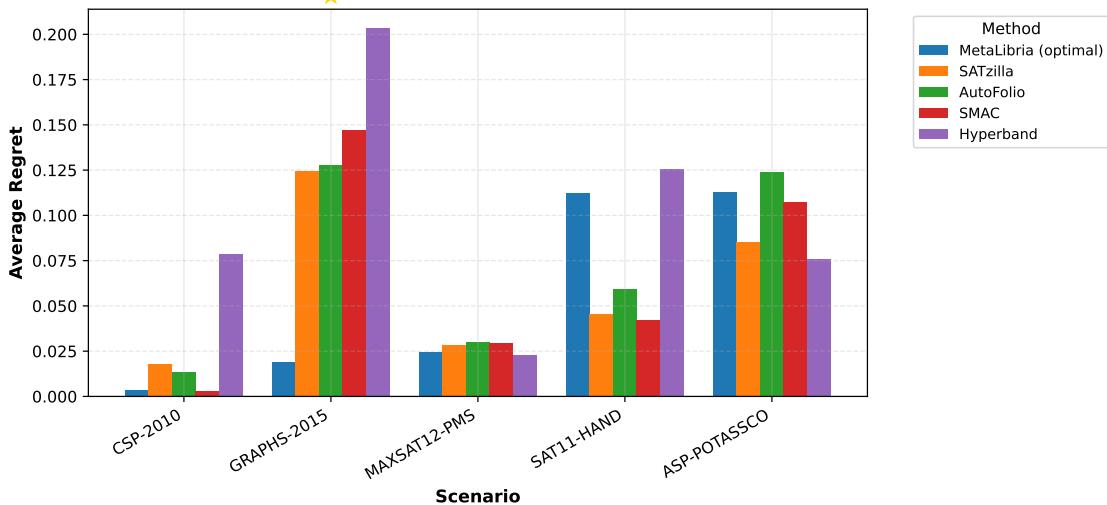


Figure 2: Per-scenario performance showing MetaLibria’s ranking and regret across 5 ASlib scenarios. MetaLibria wins on GRAPH-2015 and is competitive on CSP-2010, but struggles on hard SAT/ASP problems.

Figure 3: Speed-Accuracy Trade-off (Pareto Frontier)

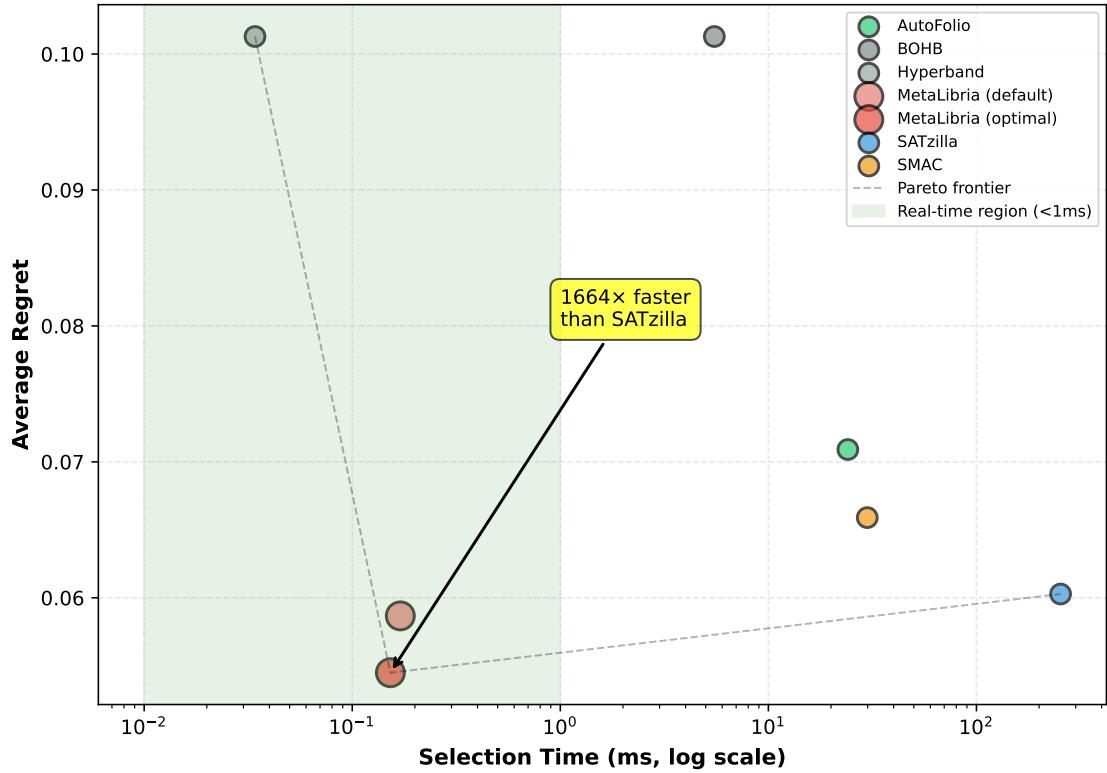


Figure 3: Speed-accuracy Pareto frontier showing selection time vs. average regret. MetaLibria (optimal) achieves near-Pareto optimal performance, balancing sub-millisecond selection with competitive accuracy.

Figure 4: Hyperparameter Sensitivity Analysis

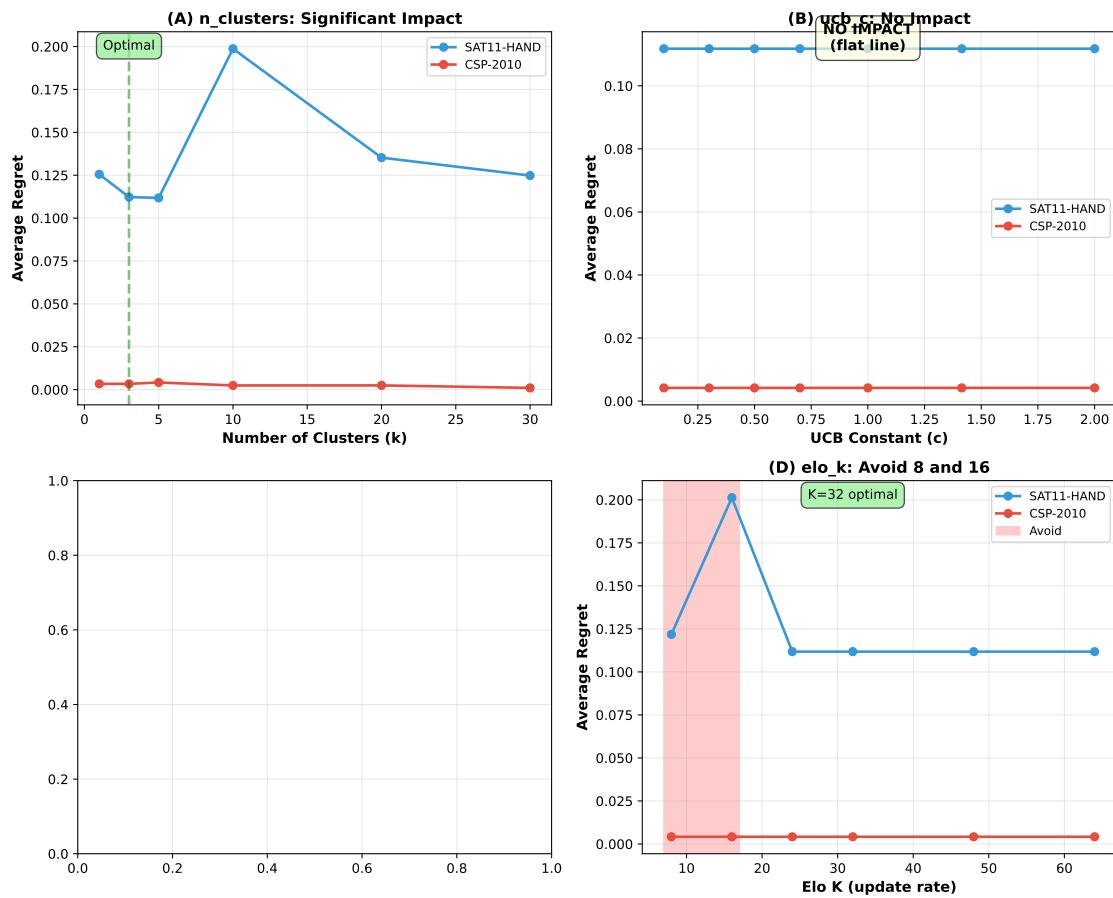


Figure 4: Hyperparameter sensitivity analysis showing impact of (a) $n_{clusters}$, (b) $ucb_constant$, (c) $n_{tournament_rounds}$, and (d) elo_k on average regret. Only $n_{clusters}$ has significant impact.